

Efficient Neutrino Oscillation Parameter Inference with Gaussian Process

Lingge Li, Nitish Nayak, Jianming Bian, Pierre Baldi

UC-Irvine

PhystatNu - 2019

Neutrino Oscillations

- ▶ Neutrinos : 2 kinds of states, each of which come in 3 types
 - ▶ Interacting, i.e what we observe → flavor states (ν_e, ν_μ, ν_τ)
 - ▶ Propagating, i.e in between observations → mass eigenstates (ν_1, ν_2, ν_3)
- ▶ Principle of superposition connects them via 3×3 unitary matrix (U_{PMNS}), i.e.

$$\begin{bmatrix} \nu_e \\ \nu_\mu \\ \nu_\tau \end{bmatrix} = U_{PMNS} \begin{bmatrix} \nu_1 \\ \nu_2 \\ \nu_3 \end{bmatrix}$$

- ▶ Via QM, neutrinos starting out as one flavor can be observed as another ("Oscillations").
- ▶ Well defined probability which depends on :
 - ▶ Energy of neutrino, E_ν and length of propagation, L
 - ▶ mass-squared splittings, $\Delta m_{32}^2, \Delta m_{21}^2$, i.e $\Delta m_{ij}^2 = m_i^2 - m_j^2$
 - ▶ U_{PMNS}
- ▶ For neutrino propagation in vacuum, the oscillation probability in all its glory:

$$P(\nu_\alpha \rightarrow \nu_\beta) = \delta_{\alpha\beta} - 4 \sum_{i>j}^3 \Re(U_{\alpha i}^* U_{\beta i} U_{\alpha j} U_{\beta i}^*) \sin^2\left(\frac{\Delta m_{ij}^2 L}{4E_\nu}\right) + 2 \sum_{i>j}^3 \Im(U_{\alpha i}^* U_{\beta i} U_{\alpha j} U_{\beta i}^*) \sin\left(\frac{\Delta m_{ij}^2 L}{4E_\nu}\right)$$

Neutrino Oscillations Contd..

- ▶ U_{PMNS} commonly parameterized as

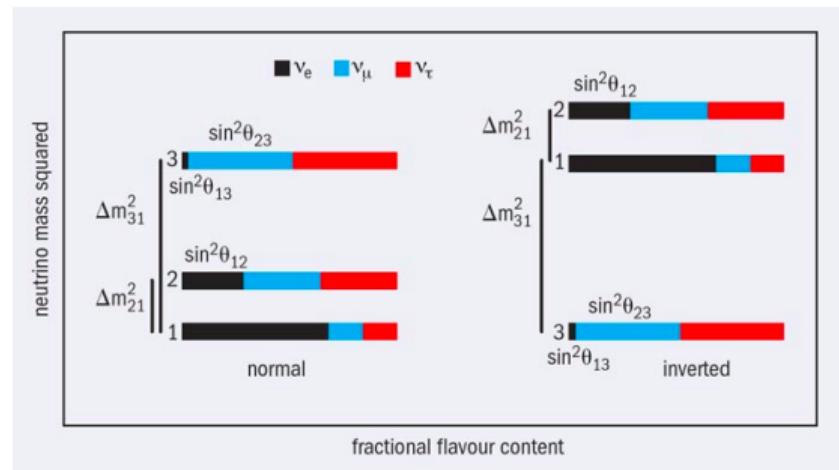
$$U_{PMNS} = \begin{bmatrix} 1 & 0 & 0 \\ 0 & \cos\theta_{23} & \sin\theta_{23} \\ 0 & -\sin\theta_{23} & \cos\theta_{23} \end{bmatrix} \begin{bmatrix} \cos\theta_{13} & 0 & \sin\theta_{13}e^{-i\delta_{CP}} \\ 0 & 1 & 0 \\ -\sin\theta_{13}e^{i\delta_{CP}} & 0 & \cos\theta_{13} \end{bmatrix} \begin{bmatrix} \cos\theta_{12} & \sin\theta_{12} & 0 \\ -\sin\theta_{12} & \cos\theta_{12} & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

- ▶ Physics program entails measuring $P(\nu_\alpha \rightarrow \nu_\beta)$ to infer U_{PMNS} and Δm_{ij}^2 parameters
- ▶ Broadly, solar experiments give handle on (21) parameters, reactor experiments for θ_{13}
- ▶ Long baseline (LBL) experiments (this talk) gives handle on (32).
 - ▶ $P(\nu_\mu \rightarrow \nu_\mu)$ sensitive to $\sin^2(2\theta_{23})$ and $|\Delta m_{32}^2|$
 - ▶ Non-zero θ_{13} opens up $P(\nu_\mu \rightarrow \nu_e)$ channel, sensitive to δ_{CP} , θ_{23} octant and $\text{sgn}(\Delta m_{32}^2)$

Physics Implications

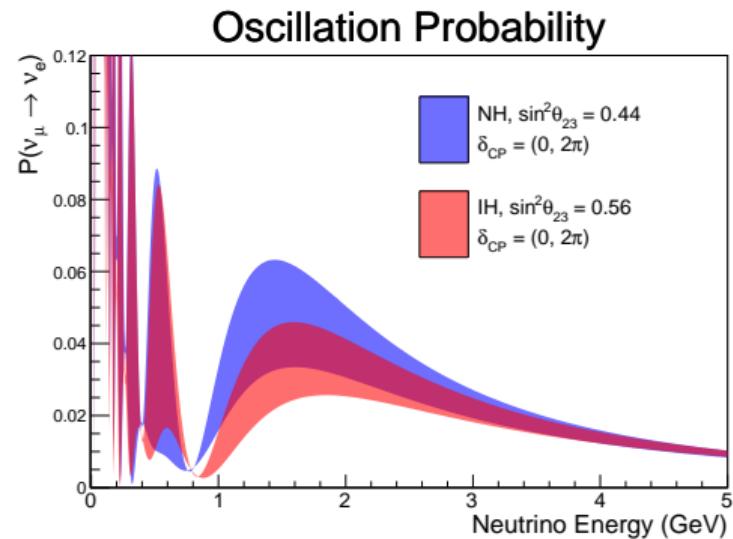
In the LBL context, we want to know if :

- ▶ $\Delta m_{32}^2 > 0$ or < 0 ? (Normal or Inverted)
 - ▶ Identifying mass hierarchy (NH or IH) has implications for neutrino mass measurements
- ▶ Octant of θ_{23} or $\theta_{23} = 45^\circ$?
- ▶ $\sin\delta_{CP} \neq 0$?
 - ▶ Lepton sector CP-violation. Gives us a clue towards explaining matter-antimatter asymmetry



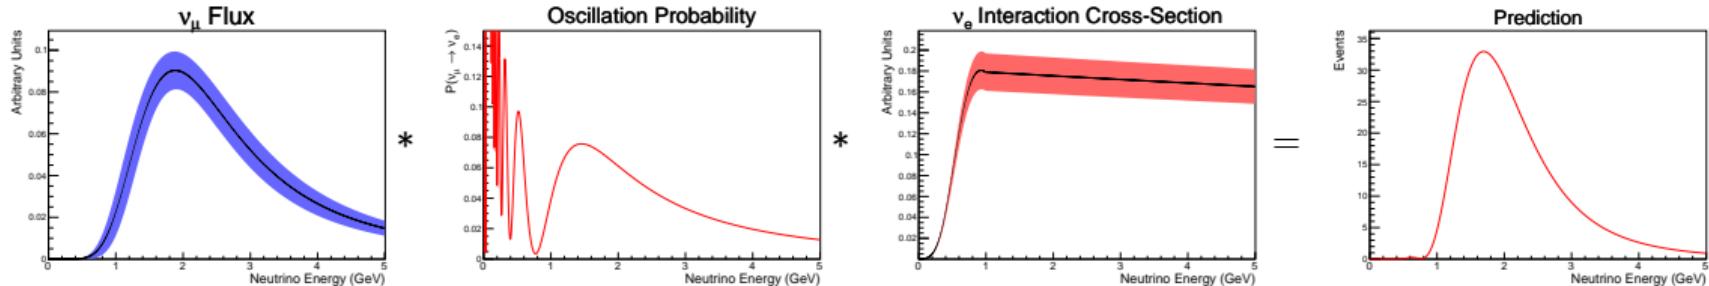
Statistical Issues

- ▶ Oscillation Parameters are typically measured via MLE using the underlying PMNS model and comparing it to observation
- ▶ However, experiments collect only a handful of statistics. $\mathcal{O}(10 - 100)$ over years of operation for the $\nu_\mu \rightarrow \nu_e$ channel
- ▶ Oscillation probabilities have complicated dependence on multiple parameters \implies difficult to delineate
- ▶ Confidence Intervals are hard to find as Likelihood ratios don't satisfy asymptotic properties.
- ▶ Let's illustrate this with a toy experiment..



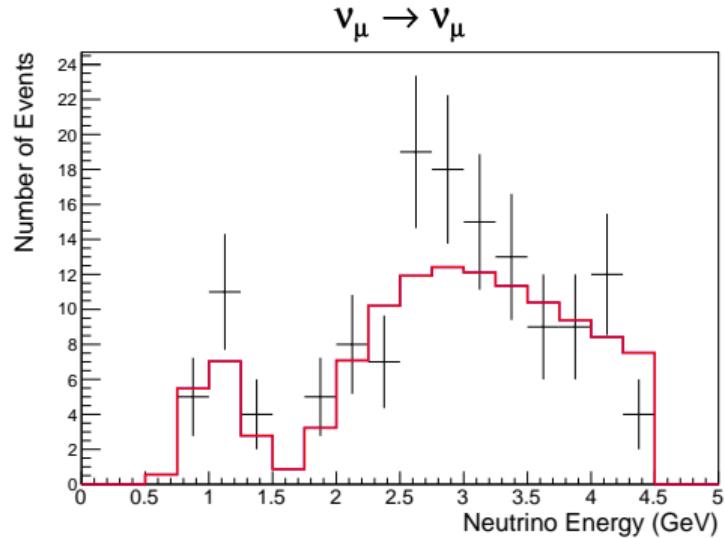
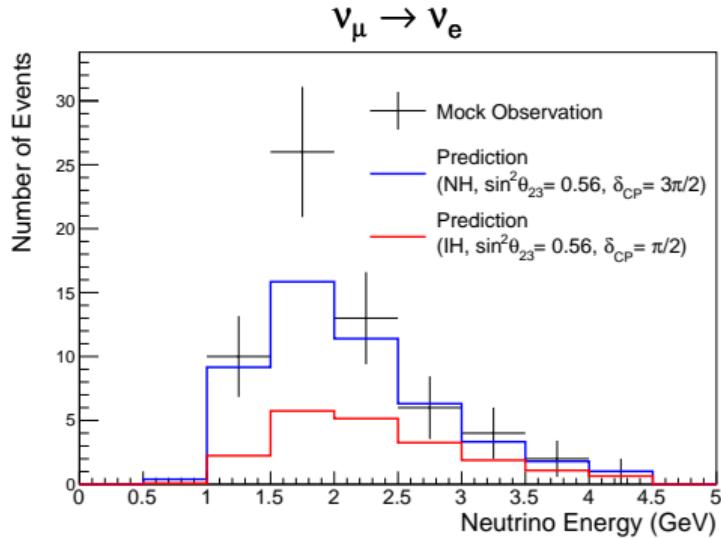
Toy Experiment

- ▶ Modelled on NOvA. Baseline, $L = 810\text{km}$ with ν_μ flux peaking at 2GeV
- ▶ $\nu_\mu \rightarrow \nu_e$ by multiplying toy shapes for flux, cross-section and oscillation probability.
- ▶ 10% normalisation errors on flux and xsec model



- ▶ $P(\nu_\mu \rightarrow \nu_e)$ using 3-flavor PMNS with MSW corrections added for matter propagation.
- ▶ Similar setup for $\nu_\mu \rightarrow \nu_\mu$ to constrain $\sin^2(2\theta_{23})$ and $|\Delta m_{32}^2|$ but with 2-flavor approximation
- ▶ $P(\nu_\mu \rightarrow \nu_\mu) \sim 1 - \sin^2(2\theta_{23})\sin^2(\Delta m_{32}^2 L / 4E)$

Toy Experiment



- Toy data (\vec{x}) from Poisson variations at some chosen oscillation parameters.
- With (θ, δ) denoting list of oscillation and nuisance (flux and xsec errors) parameters,
- Best-fit $(\hat{\theta}, \hat{\delta})$ found by minimizing negative log-likelihood over energy bins, i

$$-2 \log L(\theta, \delta) = -2 \sum_{i \in I} \log \text{Pois}(x_i; v(\theta, \delta)_i) - \sum_{i \in I} x_i + \sum_{i \in I} v(\theta, \delta)_i + \delta^2$$

Confidence Intervals

- ▶ θ_0 included in the $1 - \alpha$ confidence contour if we fail to reject the null ($\theta = \theta_0$) at α level
- ▶ Use an Inverted Likelihood Ratio Test (LRT)
- ▶ Neyman-Pearson Lemma : Likelihood Ratio (LR) is the most powerful test statistic

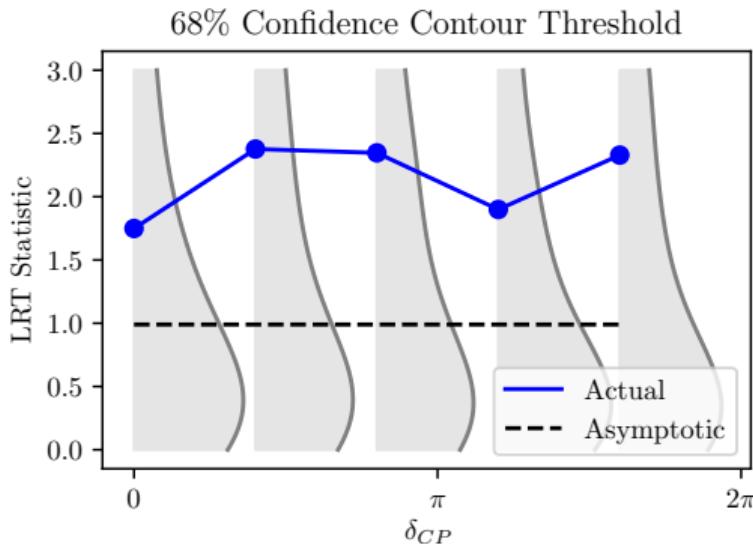
Table 38.2: Values of $\Delta\chi^2$ or $2\Delta\ln L$ corresponding to a coverage probability $1 - \alpha$ in the large data sample limit, for joint estimation of m parameters.

$(1 - \alpha)$ (%)	$m = 1$	$m = 2$	$m = 3$
68.27	1.00	2.30	3.53
90.	2.71	4.61	6.25
95.	3.84	5.99	7.82
95.45	4.00	6.18	8.03
99.	6.63	9.21	11.34
99.73	9.00	11.83	14.16

- ▶ Easy to estimate in the asymptotic case as LR is a χ^2 distribution. (Wilks Theorem)
- ▶ However, that's not the case here!
- ▶ Proceed via Unified approach (Feldman-Cousins, 1998)

From the PDG Review on Statistics

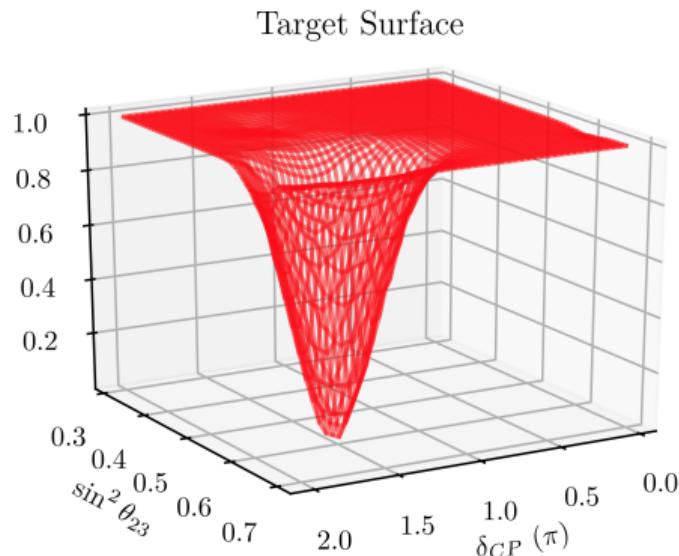
- ▶ Seminal result giving an ordering principle for confidence intervals in non-asymptotic cases
- ▶ For given θ_0 , explicitly simulate distribution of test statistic, LR via Monte-Carlo experiments at θ_0



- ▶ 68% confidence interval for δ_{CP} : All δ_{CP} values for which LR for observed data (critical value) lies within threshold
- ▶ Confidence of rejecting given $\delta_{CP} = \delta_0$ given by percentile of $\text{crit}(\delta_0)$
- ▶ Gives us the "correct" confidence interval in the frequentist sense by construction, since its essentially a grid search over the entire parameter space.

A more efficient FC

- ▶ Grid search across multi-dimensional parameter space \implies extremely intense computational demands
- ▶ It'd be nice to be able to come up with a more refined search algorithm.



- ▶ We can expect intuitively :
 - ▶ Given a point in parameter space that is rejected at high confidence, it is likely that points near it will also be rejected
 - ▶ Further, the variation in the LR percentiles ought to be smooth.
- ▶ An efficient search would therefore :
 - ▶ Learn local features in the LR percentile surface to guide the search
 - ▶ Favor simulating the LR test statistic distribution near the edge of the desired confidence contour than further out.

- ▶ Our goal is to approximate the FC percentile surface non parametrically using only a fraction of the grid points.
 - ▶ A Bayesian approach can assume a model itself to be a random variable with a certain probability distribution.
 - ▶ Training data updates your priors about the model distribution
 - ▶ Predictions for new data is a posterior distribution in model space.
 - ▶ Quantifies uncertainty in model estimates. Gets smaller with more training data
 - ▶ Can be pretty non-parametric
- ▶ Classical supervised learning → training data to get best-fit model.
 - ▶ Predictions for new data are best-guess

Gaussian Process

- ▶ Special case of Bayesian Learning. Model distribution is an extension of multivariate gaussians to function space.
- ▶ Technically, its a probability measure defined over ∞ -dim function space parameterized only by a mean function, $\mu(x)$ and a covariance function (kernel), $k(x, x')$
- ▶ We say, $f \sim \mathcal{GP}(\mu, k(\cdot, \cdot))$ if

$$\begin{pmatrix} f(x) \\ f(x') \end{pmatrix} \sim \mathcal{N}\left(\begin{bmatrix} \mu(x) \\ \mu(x') \end{bmatrix}, \begin{bmatrix} k(x, x) & k(x, x') \\ k(x, x') & k(x', x') \end{bmatrix}\right).$$

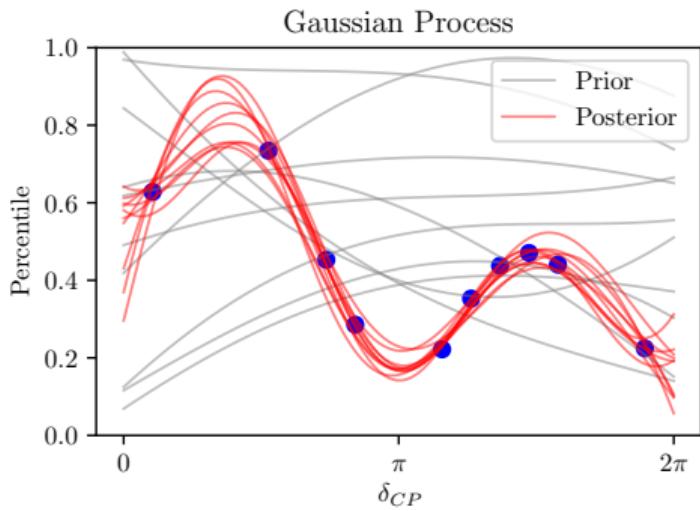
- ▶ Intuitively, we can picture each draw from a $\mathcal{GP}(\mu, k(\cdot, \cdot))$ giving us a different $f(x)$ with the average result being $\mu(x)$
- ▶ The kernel encodes the correlation between nearby points. A commonly used kernel is the radial basis function, $k(x, x') = \exp(-(x - x')^2/l^2)$
- ▶ A RBF kernel tells us that \mathcal{GP} results at nearby points are highly influenced by observations at a given point while further out, they aren't.

Why \mathcal{GP} s?

- ▶ Enormously flexible! Can basically approximate any well behaved function with an appropriate choice of the kernel.
- ▶ Predictions at new data points are computationally tractable with basic linear algebra, i.e for $\mathcal{GP}(\mathbf{0}, k(\cdot, \cdot))$:

$$f(x')|f(x) \sim \mathcal{N}\left(\frac{k(x, x')}{k(x, x)} f(x), k(x', x') - \frac{k(x, x')^2}{k(x, x)}\right)$$

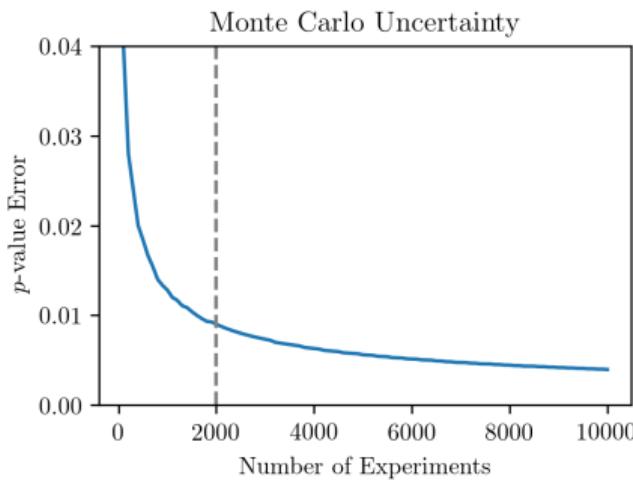
- ▶ Kernel hyperparameters can be learned via maximising the likelihood of current set of observations marginalised over the function distribution, f



\mathcal{GP} s in Literature

- ▶ \mathcal{GP} s in HEP : arXiv:1709.05681, M. Frate, K. Cranmer et al. Using \mathcal{GP} s to describe background spectra in dijet resonance searches at the LHC non-parametrically.
- ▶ Used in Astrophysics for modelling stochasticity of light yields in stars, active galactic nuclei etc
- ▶ Many other fields!

- ▶ Fitting a GP to target percentile surface for a given contour. (Stochasticity of the target surface)
- ▶ "Observation" at a given point in parameter space, θ means simulating the LRT distribution and finding the percentile of $crit(\theta)$
- ▶ Choose a RBF Kernel with an additional term incorporating variance of percentile estimate at θ .



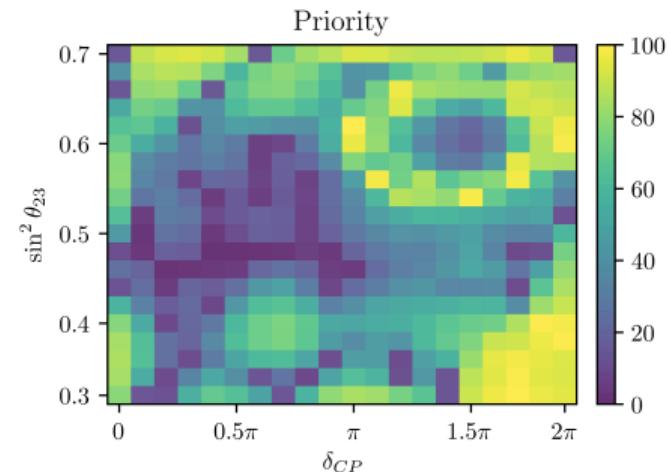
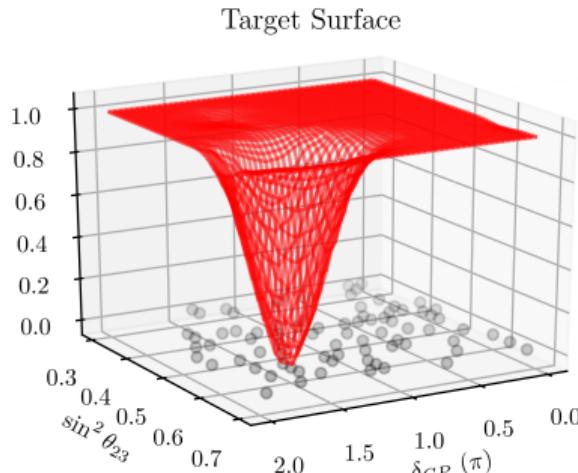
- ▶ $k(\cdot, \cdot) = k_{RBF}(\cdot, \cdot) + \sigma_p^2 I$
- ▶ The additional variance encodes the binomial error resulting from throwing finite number of experiments to simulate the LRT distribution at θ
- ▶ Allows us to incorporate varying number of experiments thrown into the CI search, reducing computational burden further.

Optimised Confidence Interval Search

- ▶ Use an acquisition function that proposes new points in θ -space to explore based on \mathcal{GP} approximated percentile surface.

$$a(\theta) = \sum_{\alpha_i} \left| \frac{\hat{q}(\theta) - \alpha_i}{\sigma_{\hat{q}(\theta)}} \right|^{-1}$$

- ▶ Here, $\hat{q}(\theta)$ is \mathcal{GP} mean, $\sigma_{\hat{q}(\theta)}$ is \mathcal{GP} std-dev, α_i is chosen to be $(0.68, 0.90)$
- ▶ $a(\theta)$ balances between exploration, i.e MC experiments at new points and exploitation, i.e reducing \mathcal{GP} error

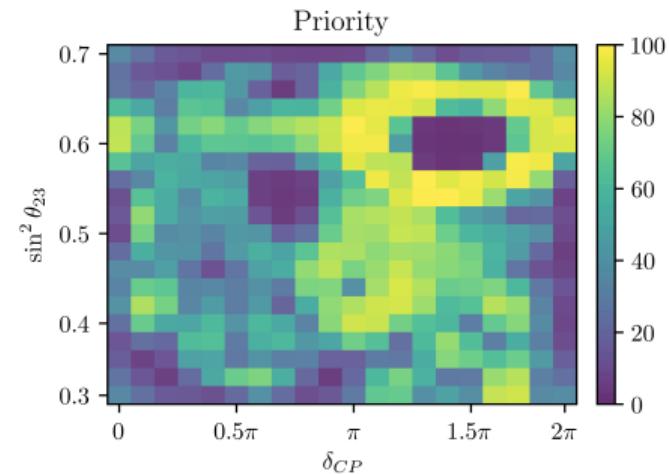
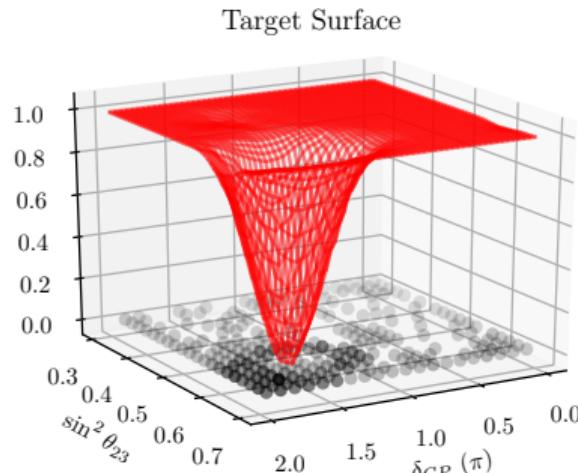


Optimised Confidence Interval Search

- ▶ Use an acquisition function that proposes new points in θ -space to explore based on \mathcal{GP} approximated percentile surface.

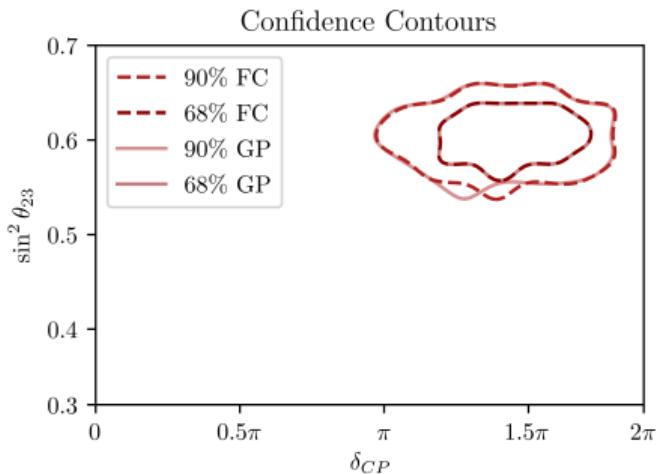
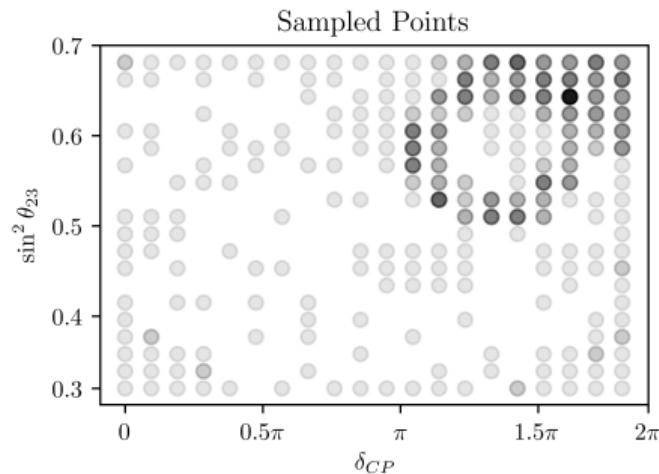
$$a(\theta) = \sum_{\alpha_i} \left| \frac{\hat{q}(\theta) - \alpha_i}{\sigma_{\hat{q}(\theta)}} \right|^{-1}$$

- ▶ Here, $\hat{q}(\theta)$ is \mathcal{GP} mean, $\sigma_{\hat{q}(\theta)}$ is \mathcal{GP} std-dev, α_i is chosen to be $(0.68, 0.90)$
- ▶ $a(\theta)$ balances between exploration, i.e MC experiments at new points and exploitation, i.e reducing \mathcal{GP} error



Results

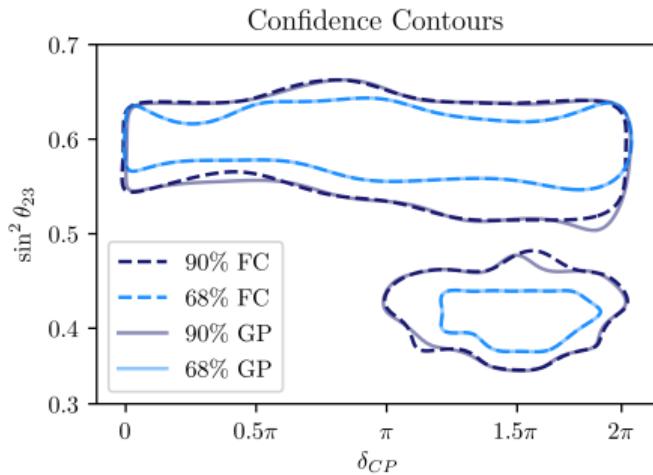
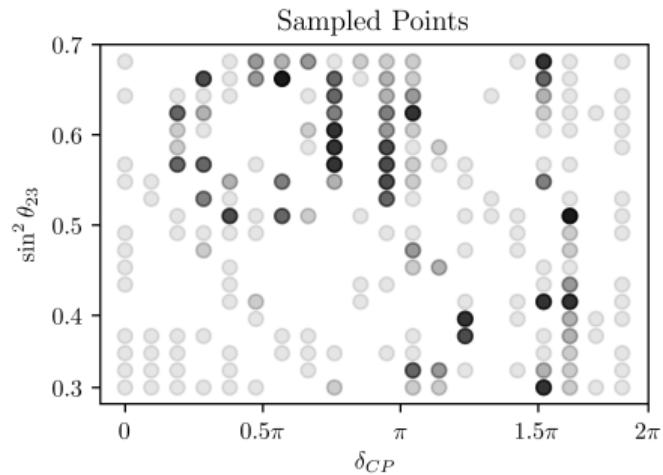
- ▶ "Real" data similar to latest best-fit estimate from NOvA. ($\sin^2 \theta_{23} = 0.56$, $\Delta m_{32}^2 = 2.44 \times 10^{-3} \text{ eV}^2$, $\delta_{CP} = 1.5\pi$)
- ▶ $\sin^2 \theta_{23} - \delta_{CP}$ 68% and 90% CI for IH after 5 iterations



- ▶ Grayscale denotes number of experiments thrown in relation to FC (2000)
- ▶ Algorithm does a good job of finding the FC contour edge!

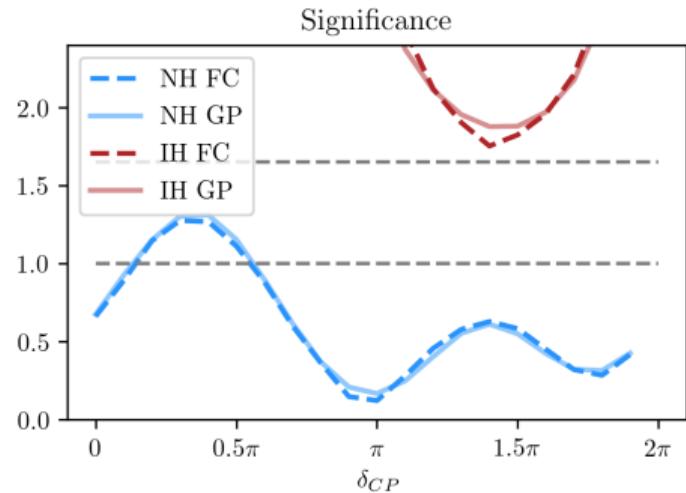
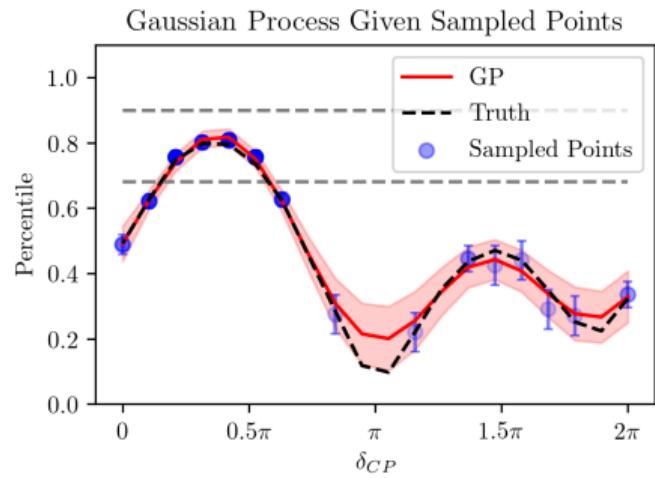
Results

- "Real" data similar to latest best-fit estimate from NOvA. ($\sin^2 \theta_{23} = 0.56$, $\Delta m_{32}^2 = 2.44 \times 10^{-3} \text{ eV}^2$, $\delta_{CP} = 1.5\pi$)
- $\sin^2 \theta_{23} - \delta_{CP}$ 68% and 90% CI for NH after 5 iterations



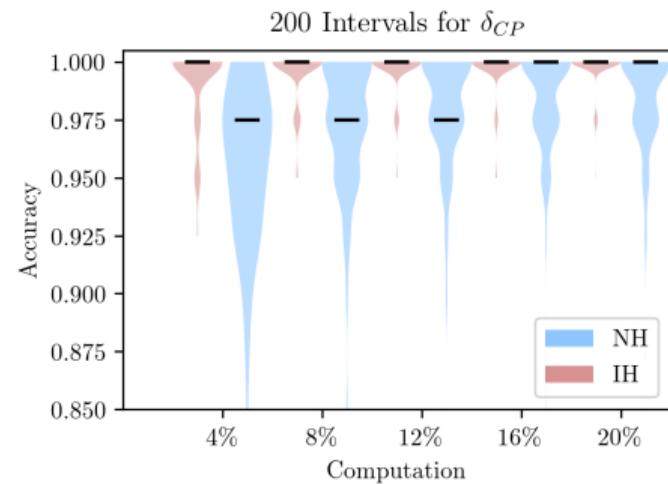
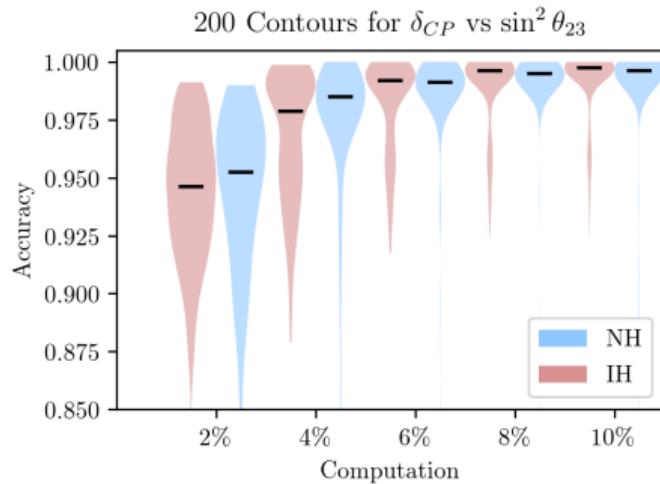
Results

- ▶ "Real" data similar to latest best-fit estimate from NOvA. ($\sin^2\theta_{23} = 0.56$, $\Delta m_{32}^2 = 2.44 \times 10^{-3} \text{ eV}^2$, $\delta_{CP} = 1.5\pi$)
- ▶ Significance of rejecting δ_{CP} only after 5 iterations. (Percentile converted to Z-score significance)



Results

- ▶ 200 different runs for "real" data at the same point as before.
- ▶ Use classification accuracy of all grid points, taking FC result as truth, to evaluate performance.
- ▶ Progress shows the search algorithm converges to the FC value $\sim 10\times$ faster for 2D case and $\sim 5\times$ for 1D case



- ▶ Median Accuracies for 1D is 100%, for 2D is $> 99.5\%$ (both NH, IH)
- ▶ Mean Accuracies for 1D is 98.5% (99.8%) for NH (IH), for 2D is $> 99\%$ (both NH, IH)

Summary and Conclusions

- ▶ Neutrino oscillation experiments provide interesting test case for estimating frequentist confidence intervals
- ▶ LBL experiments typically proceed via Feldman-Cousins
- ▶ However, simulating LRT distributions across multi-dimensional parameter space requires huge computational resources
- ▶ We've studied a Bayesian approach using Gaussian processes on a toy LBL set-up
- ▶ Helps us estimate frequentist contour edges to quite a high accuracy without having to sample the entire parameter space!
- ▶ Order of magnitude gain in computation!
- ▶ All code with illustrative notebooks here : <https://github.com/nitish-nayak/ToyNu0scCI>, maintained by Lingge (linggeli7@gmail.com) and myself (nayakb@uci.edu)

Backup

- ▶ Rasmussen and Williams has a good discussion about convergence to true functions in regression settings (typically using squared loss functions) :
<http://www.gaussianprocess.org/gpml/chapters/RW7.pdf>
- ▶ Well behaved \implies expressible as a generalised fourier series of kernel eigenfunctions
- ▶ If kernel is non-degenerate, approximation is guaranteed to converge to true function
- ▶ If degenerate, convergence towards an L_2 approximation of the true function
- ▶ Rates of convergence typically depends on mean and kernel smoothness as well as smoothness of the true function

- ▶ Hyperparameters (\mathbf{w}) learned via maximising log marginal likelihood :

$$p(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \int p(\mathbf{y}|\mathbf{X}, \mathbf{w}, \mathbf{f})p(\mathbf{f}|\mathbf{X}, \mathbf{w})d\mathbf{f}$$

- ▶ Clearly,

$$\mathbf{f}|\mathbf{X}, \mathbf{w} \sim \mathcal{N}(\mathbf{0}, K(\mathbf{X}, \mathbf{w}))$$

- ▶ Some algebra gives us :

$$-2 \log p(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \mathbf{y}^T K^{-1} \mathbf{y} + \log |K| + n \log 2\pi$$

- ▶ Minimising above equation gives us a good choice for \mathbf{w}
- ▶ $\log|K|$ acts as a penalty term for complexity and therefore reduces overfitting to data

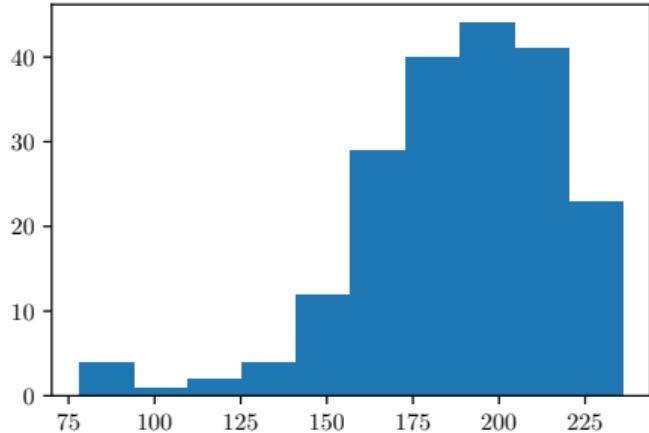
- ▶ "Gaussian" not a statement of the underlying distribution of the test statistic, which can still be heavily non-Gaussian
- ▶ Rather, "Gaussianity" for a stochastic process generating the test statistic distributions. Stochasticity mostly from finite FC grid resolution or finite number of MC experiments for simulating the test statistic distribution
- ▶ Assumption we're making for this stochasticity is that it can be parameterised by a kernel describing the relationship between the distributions at neighbouring points \implies multi-variate gaussian
- ▶ Also important to note, no real statement about FC coverage or handling of nuisance parameters. Assumes FC gives desired level of coverage
- ▶ Confidence Intervals still with frequentist interpretation
- ▶ Bayesian interpretation for "classification probability" of points in parameter space for desired confidence regions
- ▶ A good summary would be "Accelerating Frequentist CI search by estimating CI edges through Bayesian ML"

Algorithm 1 \mathcal{GP} iterative confidence contour finding

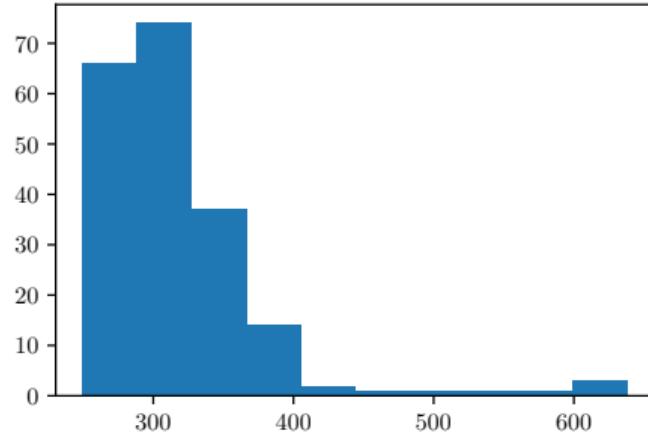
```
for each iteration  $t = 1, 2, \dots$  do
    Propose new points in parameter space  $\arg \max_{\theta} a(\theta)$ 
    for each point  $\theta'$  do
        Simulate likelihood ratio distribution
        for  $k = 1, 2, \dots$  do
            Perform a pseudo experiment
            Maximize the likelihood with respect to  $(\theta, \delta)$ 
            Maximize the likelihood with constraint  $\theta = \theta'$ 
        end for
        Obtain critical value  $c(\theta')$ 
    end for
    Update  $\mathcal{GP}$  approximation  $\hat{c}(\theta)$ 
    Update confidence contours
end for
```

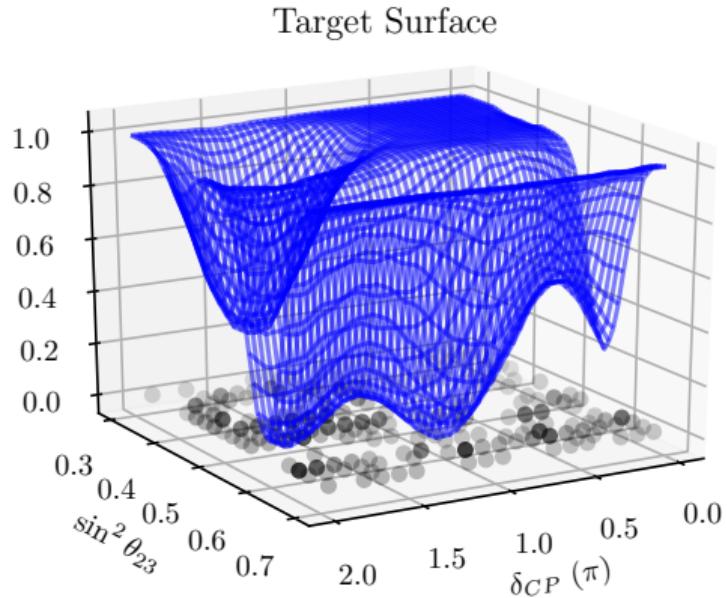
Results : NH, $\sin^2\theta_{23} - \delta_{CP}$

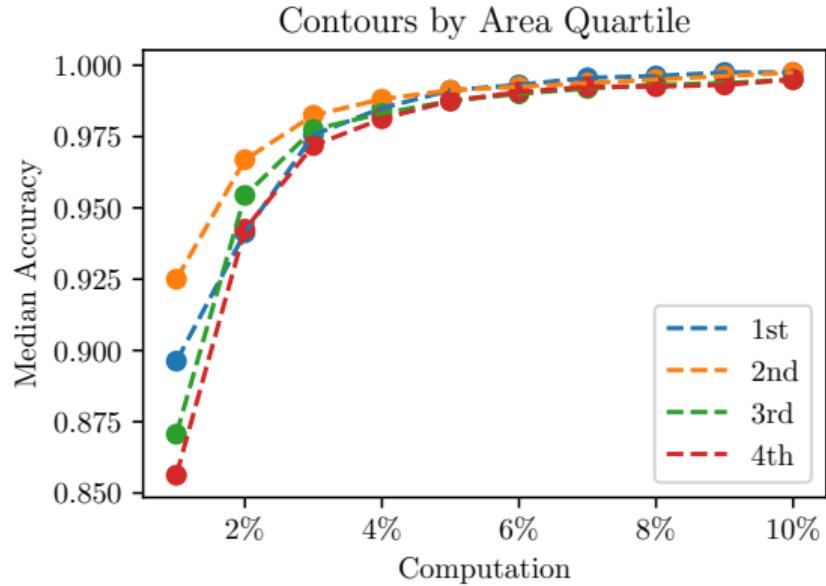
Number of Points Explored



Mean Number of Simulations







Target Surface

